

CMPT 417

Tetromino Packing

Jacob Patenaude - 301203788

Problem

Given a 4x4 grid and a number of rectangle, square, L-shaped, and T-shaped tetrominoes, determine if there exists a way to arrange the tetrominoes such that the board is filled without any overlapping pieces. E.g. there is a packing when there are four rectangles:

1	2	3	4
1	2	3	4
1	2	3	4
1	2	3	4

Chosen Language

I chose to write the specifications for the problem in the [IDP System by DTAI Research Group](#) as the logical statements involved were very close to the theoretical second-order logical formulae we have been using in class. For writing the IDP specification, I made extensive use of the [online IDE](#).

Motivation/Goals

I was interested in the tetromino packing problem as it seemed to have an obvious, but not necessarily trivial, extension –making the board size bigger. My goal was to, after solving the given problem, create a generic way of generating a satisfying model for a given number of each tetromino and to have a way to model the performance of finding that solution so that I could see how the performance changed for arbitrary inputs and even, hopefully, for arbitrary board sizes. Finally, due to the visual nature of solutions to the problem, I wanted to create a visual output of solutions to quickly and cleanly display and check solutions.

Initial Solution (Vocabulary and Structure)

To start the vocabulary, I gave each of the four blocks that would be used in the packing an id from 1 to 4. Then, to describe the placement and type of tetrominoes within the board, I used five associations:

- I. The type of a piece out of rectangle, square, T-piece, or L-piece (referred to as constants “R”, “S”, “T”, “L”).
- II. The rotation of a piece out of 0, 90, 180, or 270 degrees clockwise (referred to as constants “ROT_0”, “ROT_90”, “ROT_180”, “ROT_270”).
- III. The reflection of a piece (simply a boolean of whether or not a piece is reflected).
- IV. The location of a piece’s anchor position –I gave each of the tetrominoes a special cell to anchor a piece onto the board, given its rotation and reflection. E.g. for a rectangle, I anchored the piece by its bottommost cell (when viewing a rectangle vertically).
- V. The block that occupies the cell at a given position in the board. The actual solution to the problem is this list of block positions as it directly describes the packing which corresponds to the solution to the given inputs (if the inputs give a satisfiable structure).

These were implemented as relations:

```
BlockType(Block, Type)
Rotated(Block, Rotation)
Reflected(Block)
Located(Block, Index, Index)
Has(Index, Index, Block)
```

Where “Block” corresponds to a block’s idea, “Type” to one of the type constants, and “Rotation” to one of the rotation constants.

From here it was a simple matter of defining the input as the number of each tetromino: nR, the number of rectangles; nS, the number of squares; nT, the number of T-pieces; and nL, the number of L-pieces as the instance structure.

Lastly, I needed to write the theory which would allow IDP to derive the packing of the board from everything given above.

Initial Solution (Theory)

The simplest part of the theory was giving some constraints on the inputs. I knew that I needed to ensure that the given numbers of each tetromino corresponded to the actual number of occurrences of that block within the solution and that I needed to restrict the input relations to actually contain valid entries. That gave me the following restrictions:

```
// All blocks have a type and the number of each type is correct
!b[Block] : ?1 t[Type] : BlockType(b, t).
#{ b[Block] : BlockType(b, "R") } = nR.
#{ b[Block] : BlockType(b, "S") } = nS.
#{ b[Block] : BlockType(b, "T") } = nT.
#{ b[Block] : BlockType(b, "L") } = nL.

// All cells have a block
!x[Index], y[Index] : ?1 b[Block] : Has(x, y, b).

// All blocks have a unique location, reflection, and rotation
!b[Block] : ?1 x[XIndex], y[YIndex] : Located(b, x, y).
!b[Block] : ?1 r[Rotation] : Rotated(b, r).
```

From here, the theory became much more tedious and challenging. For each of the possible block types and rotations and for both reflected and not reflected blocks, I needed to define which cells would be occupied by a block relative to its anchor position in general terms. This gave me over a dozen statements that roughly all had the form:

```
!b[Block], x[Index], y[Index] : // For all blocks and cells
    BlockType(b, "R") & // If the block is a rectangle
    Located(b, x, y) & // And the block is anchored to the cell
    Rotated(b, "ROT_0") // And the block is not rotated
    // No mention of reflection as a rectangle reflected is the same!
=> // Then
    (y + 3 =< 3) & // The top of the rectangle is not outside the board
    Has(x, y + 0, b) & // And the rectangle occupies all the cells above it
    Has(x, y + 1, b) &
    Has(x, y + 2, b) &
    Has(x, y + 3, b).
```

And at last a solution was had! As an aside, I was initially getting incorrect solutions where the locations of blocks made no sense. I realised it was due to the provided notes having the various indexing constraints (e.g. $(y + 3 \leq 3)$ above) on the left side of the implication, but the solver was simply making the

implication false and putting the pieces wherever it wanted (within the other given constraints). Though I am curious if with a different set of constraints having the indexing occur before the implications could still give correct results, for me making this change was enough for IDP to spit out valid solutions for the given inputs.

At this stage, calling IDP for a single set of inputs took about 0.28 seconds.

Next Steps

Now that I could determine if a set of inputs corresponded to a satisfiable model and print out what that model was:

```
procedure main(){
    printmodels(modelexpand(TetrominoPacking, Packing))
} // TetrominoPacking is the theory and Packing the structure
```

I wanted to automate running the task so that I could check for packings for all possible inputs nR, nS, nT, nL for a 4 x 4 board. This amounted to finding all integer compositions of 4 of length 4 (E.g. (4, 0, 0, 0) for 4 rectangles or (1, 1, 1, 1) for one of each block type). Admittedly, I copied optimised code for doing this from [a gist](#). Now that I could iterate over all inputs, I needed a way to programmatically call IDP. To do this I added the IDP binary to my system's path (to allow calling it as `idp <filename>` from any directory) and used Python's `subprocess.call` utility to invoke that call (adding the "nowarnings" flag to keep my terminal clean from the warning that IDP was auto-completing the given structure as that was exactly the goal!). Finally, I needed a way to set the inputs repeatedly in the IDP file itself (`main.idp` in my case). To do this, I created a template file with all of the vocabulary and theory and then repeatedly copied the files contents and added the current iteration's inputs and output the results to my IDP file and that was it!

It turns out that there are 35 compositions of the integer 4 of length 4 (when 0s are included) so at about 0.3 seconds per IDP call and with some processing time, running through all of the inputs and determining the satisfying model (if any) took about 10 seconds.

Now that I had all of the valid inputs for a 4 x 4 board, I was interested in displaying them so that I could see the packing which made each board satisfiable and verify its correctness. To accomplish this, I scraped the stdout output of the `printmodels` call in my IDP main procedure by changing my use of `subprocess.call` to `subprocess.check_output` and then parsed that output to generate a JSON (JavaScript Object Notation) string containing all of the details about each iteration. This allowed me to generate a very basic HTML file with the

output (as JSON) hard-coded in under a `<script>` HTML tag. The reason I elected to do this, was that I could now generate the output using JavaScript which, in addition to being very familiar to me, meant that I didn't need to re-run my IDP script to view changes to, for example, my CSS styling.

Final Generalisations

Now that I could view all of the packings for a 4 x 4 grid, I had two final goals for generalising the problem:

1. Generate packings for all valid board sizes ranging from 2 x 2 up to some specifiable maximal size (including non-square sizes such as 2 x 4)
2. Add the final Z-piece tetromino (for all its reflections and rotations)

The first problem was overall quite straightforward, simply change the type `Index` to be two separate types `XIndex` and `YIndex`, and also allow specifying what the maximum values for those indices were (simply the number of rows/columns decremented by 1 for 0-indexing) and what the maximum block id would be. In general, for an $N \times M$ board, there are $N \times M / 4$ tetrominoes since each tetromino occupies 4 units of area and there are $N \times M$ available area units in such a board. This meant I was now generating all integer compositions for the number of blocks of length 4, rather than simply the 4 blocks I was using before.

The second problem was slightly more challenging as, in addition to supporting the new `nZ` input for the number of Z-pieces, I needed to add the theory for the cell-occupation of such a piece given its anchor location! This (like the L-piece had been) was extra annoying as a Z-piece, unlike a rectangle, is different for all 4 rotations and both reflections (reflected and not). Additionally, I now needed to generate integer compositions of length 5 to accommodate the new block type.

Now that I had ways of generalising the problem, I utilised Python's `getopt` utility to take in some command line arguments to allow for modifying, for example, the maximal board size or whether the Z-piece should be included. The final options available from the command line are copied and formatted below:

Usage:

<code>-n <numRows></code> or <code>--numrows <numRows></code>	Specifies the height of the board (Overridden by <code>--boardsize</code>)
<code>-m <numColumns></code> or <code>--numcolumns <numColumns></code>	Specifies the width of the board (Overridden by <code>--boardsize</code>)
<code>-b <maxSize></code> or <code>--boardsize <maxSize></code>	Specifies the maximum board size. Will run through all possible heights and widths up to <code><maxSize> x <maxSize></code> (Default: 4)
<code>-r</code> or <code>--reflections</code>	Include reflections (Default: False)
<code>-z</code> or <code>--ztype</code>	Include Z-pieces (Default: False)
<code>-h</code> or <code>--help</code>	Shows this help message

However, I noticed a new problem. While adding the separate x and y indices did not have a significant effect on performance, the additional constraints for the Z-piece did have a noticeable effect, even when the Z-pieces were not included. At the 4 x 4 board size level, the average time spent making an IDP call increased from about 0.28 seconds to about 0.36 seconds. I realised that even though the Z-piece constraints did not need to be fully evaluated when Z-pieces were not included (as they obviously only apply to Z-pieces), simply enumerating them all was resulting in an unnecessarily long grounding that was impacting performance. By modifying my templating to comment out the Z-piece constraints when they were not needed, the previous performance was restored.

Performance Improvements

I am confident that there exists a modification of my every block and cell theories that would improve performance; however, I was unable to find it. One thought I had was to assert some facts about rectangles, squares, and T-pieces to take advantage of their symmetries (e.g. asserting that all squares should not be rotated or reflected); however, the effect was actually slightly unfavourable. Additionally, I tried combining the for every block and cell theories of each block type into one larger theory; however, this also had a slightly unfavourable effect. It is unclear whether any combination of these changes would have a beneficial effect at larger board sizes. Finally, I tried `print(onemodel(Theory, Structure))`; however, that also slightly increased runtimes. I do suspect that there are ways I could limit which combinations of tetrominoes I need to check (e.g. 3 rectangles

and an L-piece definitely can't pack a 4 x 4 board because the L-piece adds exactly one piece in the x direction that is not countered by any of the rectangles meaning that there is a single cell that cannot be matched; however, I did not get a chance to spend any time proving any such constraints (such as the incorrect restraint "nT + nL is even" which is countered by the input nR=1 nS=0 nT=2 nL=1).

Predicting Runtimes and Conclusions

At this stage I felt that there was little I could do to change or improve my code without finding some hidden constraint or having an epiphany about the underlying logic or even IDP itself, so I moved onto the final stage: profiling and predicting runtime performance. To start with, I ran through all combinations of boards from 2 x 2 up to 4 x 4 with each combination of reflections allowed or disallowed and Z-pieces included or not included. Please note that I ran the program on Ubuntu with an Intel(R) Core(TM) i5-4670K CPU clocked at 3.40GHz. These results are shown below:

No reflections or Z-pieces:

Board Size:	2x2	2x4	3x4	4x4
Average IDP call:	0.179s	0.191s	0.179s	0.297s
Number packable:	1 out of 4	3 out of 10	5 out of 20	9 out of 35
Total time taken:	0.716s	1.912s	3.582s	10.390s

Only reflections:

Board Size:	2x2	2x4	3x4	4x4
Average IDP call:	0.192s	0.197s	0.234s	0.381s
Number packable:	1 out of 4	3 out of 10	5 out of 20	9 out of 35
Total time taken:	0.768s	1.974s	4.686s	13.351s

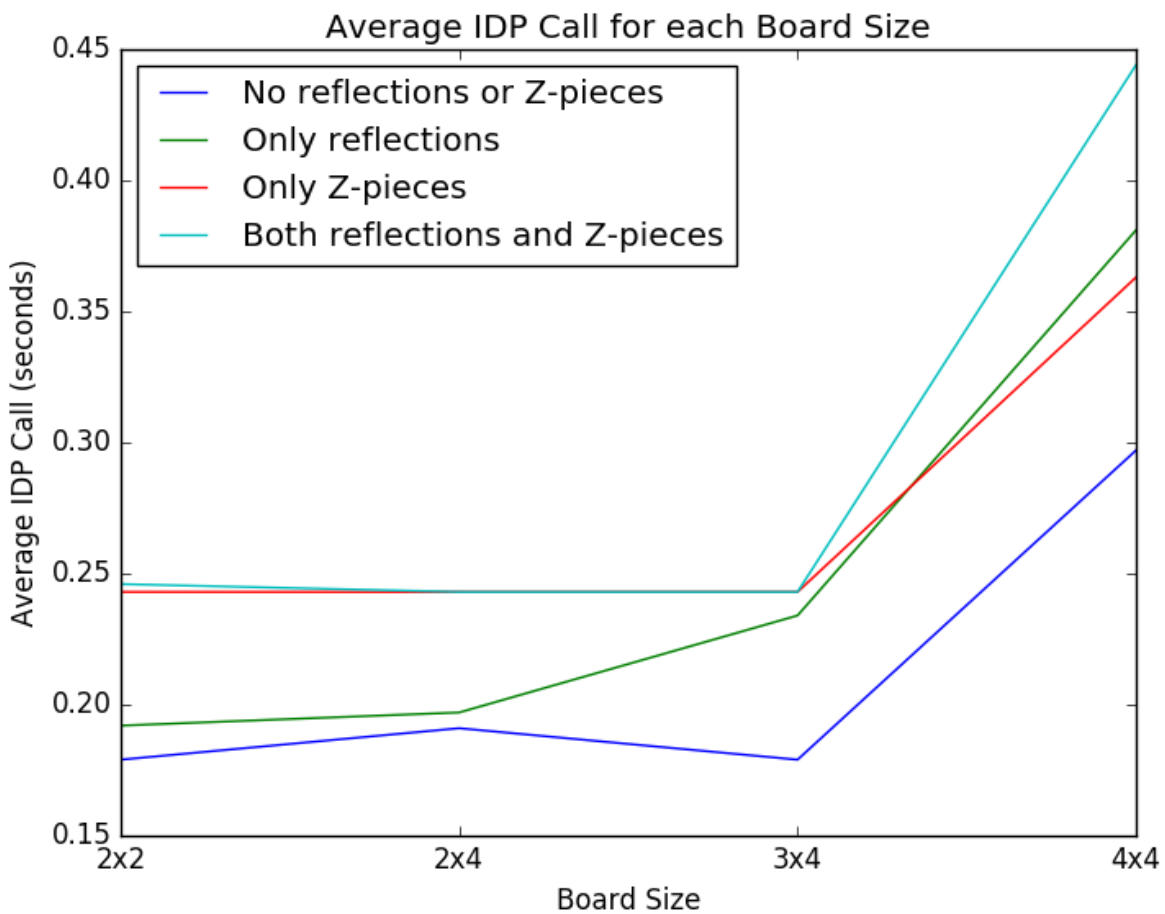
Only Z-pieces:

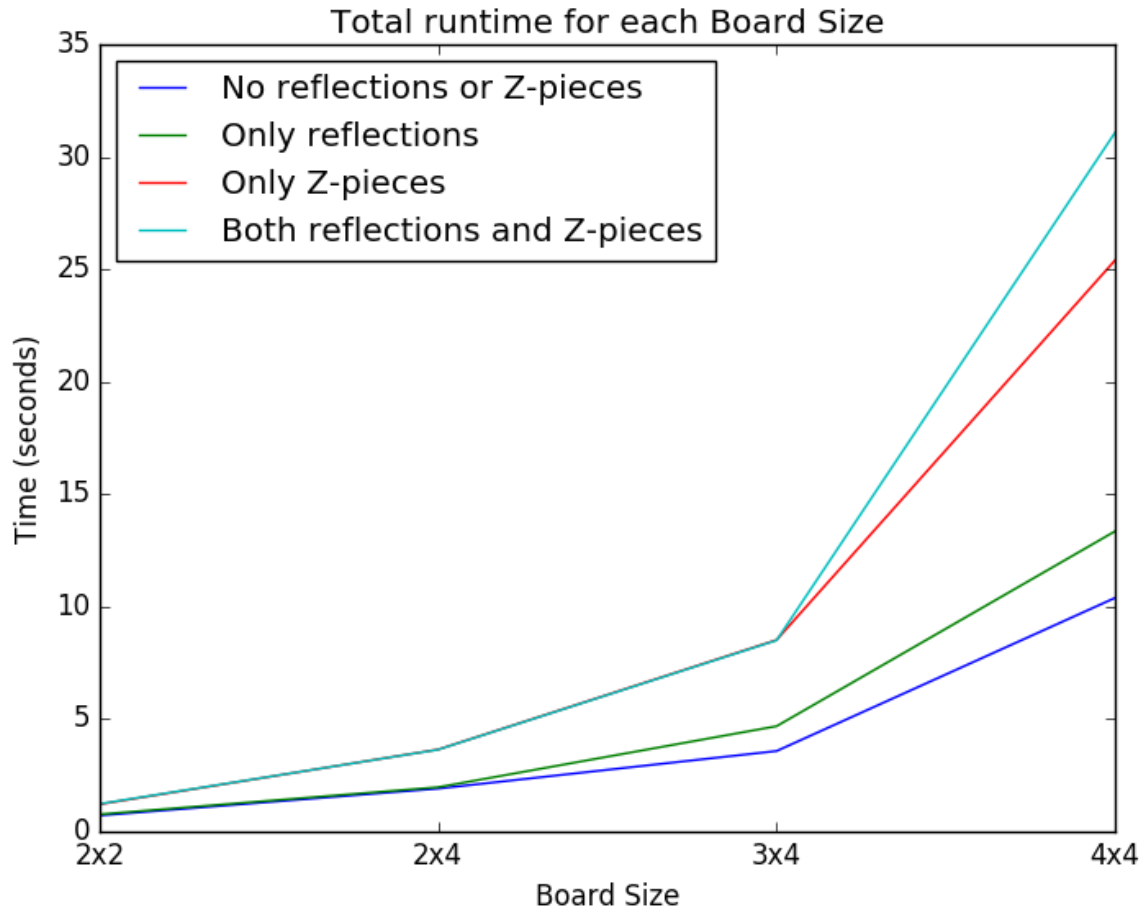
Board Size:	2x2	2x4	3x4	4x4
Average IDP call:	0.243s	0.243s	0.243s	0.363s
Number packable:	1 out of 5	3 out of 15	6 out of 35	10 out of 70
Total time taken:	1.216s	3.651s	8.521s	25.398

Both reflections and Z-pieces:

Board Size:	2x2	2x4	3x4	4x4
Average IDP call:	0.246s	0.243s	0.243s	0.444s
Number packable:	1 out of 5	3 out of 15	6 out of 35	12 out of 70
Total time taken:	1.232s	3.644s	8.503s	31.075s

Then, to help visualise this data, I plotted the time taken using the PyPlot utility available under the Matplotlib library. Please note that while code is provided to generate these plots, running it will require you to have the Matplotlib library available to your python install.





This data roughly followed what I was expecting. Interestingly, after observing the output I noted that for all boards up to 4 x 4 in size, reflections only mattered in the case of exactly one board ($nR=0$, $nS=0$, $nT=0$, $nL=2$, $nZ=2$) which helps to explain the slight jump at the 4 x 4 size between including reflections or not when Z-pieces are included. The difference between including Z-pieces and not is clear as I have already mentioned that including the theory for the Z-pieces comes at a penalty of a greatly increased grounding size due to needing a constraint on every cell and block for all 8 orientations of the Z-piece. Note that the grounding when not including Z-pieces has about 2200 terms and the grounding when including Z-pieces has about 2500 terms. Both of these values were found by observing the output of IDP's printgrounding command while reflections were included. Finally, I conclude that the performance will increase sharply still for additional instances with larger board sizes as the shapes and nature of the problem cause me to conclude the growth is heavily exponential. An additional consideration is that while for a 4 x 4 grid and smaller, the number of possible positions for each tetromino is limited by the indices of the board, for anything larger than 4 x 4 there will be many more possible positions for each tetromino. E.g. the number of possible positions for a rectangle more than doubles from a 4 x 4 grid to a 5 x 5 grid. To test my hypothesis I will run my program for as long as possible with Z-pieces and reflections included and attach the output below.